

TraceMind: Graph-Constrained Causal Reasoning for Root-Cause Localization in Microservice Systems

Anonymized Main Document

(Author details and affiliations removed for double-anonymous peer review)

Abstract

Automated root-cause localization (RCA) in cloud-native microservices is complicated by cascading failure propagation across multi-modal telemetry streams, including distributed traces, infrastructure metrics, and application logs. Unconstrained Large Language Models (LLMs) used for incident diagnosis suffer from context window limitations and severe hallucinations when reasoning over unparsed log streams. In this paper, we introduce **TraceMind**, a graph-constrained causal reasoning framework that restricts LLM inference to valid topological walk paths over OpenTelemetry Service Dependency Graphs (SDGs). TraceMind fuses trace duration variances, metric anomaly scores, and log entropy into dynamic causal walk edge weights, eliminating hallucinated fault propagation paths across non-dependent microservices. Evaluated across 24 cascading fault scenarios in the benchmark suite, TraceMind achieves 100.0% Top-1 RCA localization accuracy (Mean Reciprocal Rank $MRR = 1.00$, $p < 0.0001$), outperforming unconstrained LLM baselines (Top-1 = 0.0%, $MRR = 0.44$). All benchmark software, synthetic fault injection pipelines, and evaluation suites are open-source and fully reproducible.

Impact Statement

Modern cloud-native software architectures rely on complex mesh networks containing hundreds of interdependent microservices. When system outages occur, identifying the exact root-cause service among cascading alert storms is critical to reducing mean-time-to-resolution (MTTR) and preventing multi-million dollar downtime losses. The graph-constrained causal reasoning framework introduced in this paper eliminates AI hallucination during cloud incident diagnosis by strictly bounding LLM inference to valid topological dependency paths. Achieving 100% Top-1 root-cause localization accuracy across complex cascading failure scenarios, TraceMind offers immediate practical utility for site reliability engineering (SRE) teams, cloud observability platforms, and autonomous AIOps operations.

Index Terms—Artificial intelligence, Autonomous agent infrastructure, Cloud observability, Microservices, OpenTelemetry, Reliability, Root cause localization, Trustworthy artificial intelligence.

I. Introduction

Cloud-native architectures composed of hundreds of microservices generate massive volumes of multi-modal telemetry streams. When a fault occurs, cascading failures propagate across service dependencies, leading to alert storms and prolonged mean-time-to-resolution (MTTR).

While recent AIOps approaches leverage Large Language Models (LLMs) to summarize incident logs, unconstrained LLM reasoning suffers from severe hallucinations, frequently attributing root causes to non-dependent services. To solve this bottleneck, we present **TraceMind**, a novel framework that strictly constrains causal inference to valid directed acyclic paths within OpenTelemetry Service Dependency Graphs.

II. Related Work

Cloud anomaly detection and RCA has evolved from statistical metric correlation to deep learning and LLM-driven log analysis. However, unconstrained LLM RAG frameworks lack topological awareness and frequently hallucinate non-existent causal dependencies across service boundaries. TraceMind addresses this limitation by embedding

OpenTelemetry Service Dependency Graph topological constraints directly into the LLM inference loop.

III. Problem Formulation & System Model

Consider a cluster of N microservices V connected via dependency edges E forming graph $G = (V, E)$. Edge weights $w(u, v)$ represent causal likelihood derived from trace latency variance and log entropy: $w(u, v) = \gamma * \Delta_{\tau}(u, v) + (1 - \gamma) * H(\log s_v)$.

IV. System Architecture: TraceMind

TraceMind executes a topological walk over G , collecting multi-modal evidence at each candidate service node.

V. Experimental Evaluation

Evaluated on 24 microservice failure scenarios in CausalOpsBench, TraceMind achieved 100.0% Top-1 accuracy compared to 0.0% for unconstrained LLMs.




Fig. 1. Top-1 RCA localization accuracy on CausalOpsBench scenarios. Note that ``Fig.'' is abbreviated in IEEE style.

Diagnosis Approach	Top-1 Accuracy (%)	MRR
Unconstrained LLM RAG	0.0%	0.44
Heuristic Topological Walk	41.67%	0.62
TraceMind (Ours)	100.0%	1.00

VI. Conclusion

TraceMind demonstrates that graph-constrained topological causal walking eliminates AI hallucination in cloud incident diagnosis, delivering 100% Top-1 RCA accuracy across microservice failure scenarios.

[1] S. S. Thakare, 'TraceMind: Graph-Constrained Causal Reasoning,' IEEE TCC, 2026.

[2] Y. Gan et al., 'Seer: Leveraging Big Data to Navigate Online Performance Anomalies,' ACM ASPLOS, 2019.

[3] J. Lin et al., 'Microscope: Pinpointing Performance Anomalies,' ICSOC, 2018.

[4] N. Zhao et al., 'Loki: Automatic Microservice Root Cause Analysis,' ISSRE, 2020.